

```
In [ ]: # MNIST dataset
import torchvision
import torch
import torchvision.transforms as transforms
import numpy as np
```

```
In [ ]: import torch.nn as nn

torch.manual_seed(42)
```

```
Out[ ]: <torch._C.Generator at 0x248c6a5ea70>
```

```
In [ ]: class ConvNet(nn.Module):
    def __init__(self):
        super(ConvNet, self).__init__()
        # 1 x 28 x 28
        self.layer1 = torch.nn.Sequential(
            torch.nn.Conv2d(1, 32, kernel_size=5, stride=2, padding=0), # 32 x 12 x 12
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=2, stride=2)) # 32 x 6 x 6
        self.layer2 = torch.nn.Sequential(
            torch.nn.Conv2d(32, 16, kernel_size=5, stride=1, padding=0), # 64 x 2 x 2
            torch.nn.ReLU(),
            torch.nn.MaxPool2d(kernel_size=2, stride=2)) # 16 x 1 x 1
        self.fc1 = torch.nn.Linear(16, 10)

    def forward(self, x):
        out = self.layer1(x)
        out = self.layer2(out)
        out = out.reshape(out.size(0), -1)
        out = self.fc1(out)
        return out
```

```
In [ ]: dataset_test = torchvision.datasets.MNIST(root='./datasets', train=False, transform=
test_loader = torch.utils.data.DataLoader(dataset=dataset_test, batch_size=10000, s

dataset_train = torchvision.datasets.MNIST(root='./datasets', train=True, transform=
train_loader = torch.utils.data.DataLoader(dataset=dataset_train, batch_size=100, s
```

```
In [ ]: dataset_train[0][0].shape, len(dataset_train)
```

```
Out[ ]: (torch.Size([1, 28, 28]), 60000)
```

```
In [ ]: class Config:
    lr = 0.001
```

```
In [ ]: config = Config()
```

```
In [ ]: import lightning as L
```

```

class LightningNN(L.LightningModule):
    def __init__(self, model, config):
        super().__init__()
        self.model = model
        self.loss = torch.nn.CrossEntropyLoss()
        self.config = config

    def forward(self, x):
        return self.model(x)

    def configure_optimizers(self):

        return torch.optim.Adam(self.model.parameters(), lr=self.config.lr)

    def training_step(self, batch, batch_idx):
        x, y = batch
        y_hat = self.model(x)
        loss = self.loss(y_hat, y)
        return loss

    def predict_step(self, batch, batch_idx, dataloader_idx=0):
        x, y = batch
        y_hat = self.model(x)
        y_pred = torch.argmax(y_hat, dim=-1)
        return y_pred

```

```

In [ ]: # model = AlexNet()
        model = ConvNet()

```

```

In [ ]: lightning_model = LightningNN(model, config)

```

```

In [ ]: trainer = L.Trainer(limit_train_batches=100, max_epochs=20)
        trainer.fit(model=lightning_model, train_dataloaders=train_loader)

```

```

GPU available: True (cuda), used: False
TPU available: False, using: 0 TPU cores
IPU available: False, using: 0 IPUs
HPU available: False, using: 0 HPUs

```

|   | Name  | Type             | Params |
|---|-------|------------------|--------|
| 0 | model | ConvNet          | 13.8 K |
| 1 | loss  | CrossEntropyLoss | 0      |

```

13.8 K    Trainable params
0         Non-trainable params
13.8 K    Total params
0.055     Total estimated model params size (MB)

```

```
Epoch 19: 100%|██████████| 100/100 [00:01<00:00, 83.15it/s, loss=0.0482, v_num=19]
```

```
`Trainer.fit` stopped: `max_epochs=20` reached.
```

```
Epoch 19: 100%|██████████| 100/100 [00:01<00:00, 82.94it/s, loss=0.0482, v_num=19]
```

```
In [ ]: def test(trainer, lightning_model, dataloader):
        y_pred = trainer.predict(model=lightning_model, dataloaders=dataloader)
        y_pred = torch.cat([y for y in y_pred], dim=0).view(-1)
        y_gt = torch.cat([y for x, y in dataloader], dim=0)
        acc = (y_pred == y_gt).sum().item() / y_gt.size(0)
        return acc
```

```
In [ ]: test(trainer, lightning_model, train_loader), test(trainer, lightning_model, test_l
```

c:\Users\suzy\miniconda3\envs\abbasi\lib\site-packages\lightning\pytorch\trainer\connectors\data\_connector.py:229: PossibleUserWarning: The dataloader, predict\_dataloader 0, does not have many workers which may be a bottleneck. Consider increasing the value of the `num\_workers` argument` (try 20 which is the number of cpus on this machine) in the `DataLoader` init to improve performance.

```
category=PossibleUserWarning,
Predicting DataLoader 0: 100%|██████████| 600/600 [00:03<00:00, 160.27it/s]
Predicting DataLoader 0: 100%|██████████| 1/1 [00:00<00:00, 5.86it/s]
```

```
Out[ ]: (0.9701166666666666, 0.969)
```

Number of bits per image:  $28 * 28 * 8 = 6272$  (number of features)

Number of images: 60000

Total number of bits:  $6272 * 60000 = 376320000$

Architecture 1:

- Layer 1: input channels = 1, output channels = 32, kernel size = 5, stride = 2
- Layer 2 (MaxPool): kernel size = 2, stride = 2 (4 : 1)
- Layer 3: input channels = 32, output channels = 64, kernel size = 5, stride = 1
- Layer 4 (MaxPool): kernel size = 2, stride = 2 (4 : 1)
- Layer 5: input channels = 64, output channels = 100
- Layer 6: input channels = 100, output channels = 10

MEC =  $4 * 3 * 28 * 28 / 64 + (64 + 1) * 100 + (100 + 1) * 10 = 3136 + 6500 + 1010 = 10646$

Accuracy = 0.9823

Architecture 2:

- Layer 1: input channels = 1, output channels = 32, kernel size = 5, stride = 2
- Layer 2 (MaxPool): kernel size = 2, stride = 2 (4 : 1)
- Layer 3: input channels = 32, output channels = 64, kernel size = 5, stride = 1
- Layer 4 (MaxPool): kernel size = 2, stride = 2 (4 : 1)
- Layer 5: input channels = 64, output channels = 10

MEC =  $4 * 3 * 28 * 28 / 64 + (64 + 1) * 10 = 3136 + 650 = 3786$

Accuracy = 0.97575

### Architecture 3:

- Layer 1: input channels = 1, output channels = 32, kernel size = 5, stride = 2 ()
- Layer 2 (MaxPool): kernel size = 2, stride = 2 (4 : 1)
- Layer 3: input channels = 32, output channels = 16, kernel size = 5, stride = 1
- Layer 4 (MaxPool): kernel size = 2, stride = 2 (4 : 1)
- Layer 5: input channels = 16, output channels = 10

$$\text{MEC} = 4 \times 3 \times 28 \times 28 / 16 + (16 + 1) \times 10 = 3136 + 170 = 3306$$

$$\text{Accuracy} = 0.9701$$

```
In [ ]: from mec import calculate_mec
```

```
In [ ]: X = dataset_train.data.numpy()
X = X.reshape(X.shape[0], -1) # flatten
y = dataset_train.targets.numpy()
```

```
In [ ]: X.shape, y.shape
```

```
Out[ ]: ((60000, 784), (60000,))
```

```
In [ ]: # Upper bound MEC 3 Layer NN
calculate_mec(X, y, task="classification")
```

```
Out[ ]: 3681.371774841067
```

My Method 3 has a MEC of 3306 which is smaller than 3681 (MEC required for a 3 layer NN). This suggests that the use of CNNs is effective.